

**BHASKARACHARYA NATIONAL INSTITUTE FOR
SPACE APPLICATIONS AND GEO-INFORMATICS**
WEEKLY PROGRESS REPORT (18/05/2026 – 25/05/2026)

WEEK 1

Project Name	OSINT Collection from Social Media Network
PROJECT DESCRIPTION :	Open source intelligence collection from social media network and capturing the trend from the social media platform X (formerly Twitter). The project involves automated scraping, entity extraction, graph database construction using Neo4j, and OSINT analysis of India-based political and media handles.
GROUP MEMBER :	Rishav Pal, Sayantan Mandal, Shivam Kumar
GROUP ID :	2026G268
GROUP GUIDE :	Dr. JHUMMARWALA ABDUL
GROUP COORDINATOR :	Dr. JHUMMARWALA ABDUL
COLLEGE GUIDE NAME :	Dr. SHWETA SAHARAN
COLLEGE GUIDE No. :	9509676076
COLLEGE GUIDE EMAIL. :	SHWETA.SAHARAN@GSV.AC.IN
COLLEGE NAME :	GATI SHAKTI VISHWAVIDYALAYA

WEEK 1 — DAILY WORK LOG

Day 1–2 | 18 May – 19 May 2026 | Foundation & OSINT Fundamentals

The first two days were dedicated to building the foundational knowledge required for the project. The team focused on understanding the core concepts of Open Source Intelligence (OSINT), graph databases, and the tools that would be used throughout the project week.

Topics Studied:

- Graph Databases — Concepts of nodes, relationships, properties, and graph traversal. Understanding how graph databases differ from relational databases and why they are ideal for social network analysis.
- Neo4j & Cypher Language — Installed Neo4j Desktop, explored the browser interface, practised writing Cypher queries including MATCH, MERGE, CREATE, WHERE, RETURN, and ORDER BY clauses. Studied constraint creation and index management for performance optimisation.
- Neo4j Bloom — Explored Bloom as a visual graph exploration tool. Learnt how to apply styles, filters, and search templates to visually analyse graph data without writing Cypher.
- OSINT Methodology — Studied the OSINT lifecycle: planning, collection, processing, analysis, and dissemination. Understood the ethical and legal boundaries of open-source data collection.
- Google Dorking — Practised advanced Google search operators (site:, filetype:, inurl:, intitle:, intext:, cache:, link:) to extract publicly available information about a target individual. Applied dorks to find exposed documents, social profiles, email patterns, and organisational affiliations.
- Practical OSINT on a Target Person — Performed a hands-on OSINT exercise on a designated target profile. Used Google Dorks, LinkedIn, publicly available databases, and social media search to compile a structured intelligence profile including name, designation, contact hints, institutional affiliations, and digital footprint mapping.

Day 3 | 20 May 2026 | Project Architecture & Environment Setup

With foundational knowledge established, Day 3 shifted focus to defining the project architecture and setting up the development environment for the linkmap OSINT pipeline.

- Designed the modular project folder structure: linkmap/ with sub-modules for scraper/, storage/, graph/, config/, utils/, and data/ directories.
- Set up Python 3.11 virtual environment and installed all dependencies: playwright, pydantic, loguru, tenacity, tqdm, ujson, aiofiles, and the neo4j Python driver.
- Ran playwright install chromium to configure the headless browser environment.
- Defined the data schema using Pydantic v2 BaseModel — UserProfile and TweetData models with full validation, JSON serialisation, and soft validators for edge-case tweet types.
- Created the custom logger module (utils/logger.py) with dual-handler output: colourised console and rotating file logs with 10MB rotation, 7-day retention, and zip compression.
- Configured config/settings.json with all timing, path, and scraping parameters. Finalised config/handles.json with 20 India-based political and media handles.
- Created config/topic_map.json mapping keywords to 12 political/social topic categories relevant to the Indian media ecosystem.

Day 4 | 21 May 2026 | Scraper Development & Anti-Ban Implementation

Day 4 was focused on building the core scraping engine using Playwright, along with session management, anti-ban rate limiting, and entity extraction.

- Built scraper/browser.py — persistent Chromium browser context with randomised viewport (1200–1600px), rotating user-agent strings, disabled webdriver property, and stealth JS patches to avoid bot detection.
- Built scraper/auth.py — full X/Twitter login automation with cookie persistence. Login credentials loaded from environment variables. Exponential backoff via tenacity (max 3 retries). Session cookies saved per-account to avoid re-login on subsequent runs.
- Built scraper/rate_limiter.py — human-scroll simulation with random pixel amounts (300–700px), random between-action delays loaded from settings.json, 30–90 second cooldowns between handles, and a session health check that triggers a pause after 500 scraped tweets.

- Built scraper/tweet_scraper.py — per-handle profile scraper (display name, bio, follower count, verified status) and timeline tweet collector. Used data-testid selectors for resilience. Implemented show-more click expansion, timestamp extraction, and engagement count parsing (K/M suffixes).
- Built scraper/parser.py — entity extraction pipeline: hashtag extraction (Unicode-aware regex), @mention extraction, link extraction with t.co expansion via anchor href mapping, URL newline collapse fix, plain URL fallback, and keyword-based topic inference from topic_map.json.
- Integrated parser.enrich() into tweet_scraper._extract_tweet() so every tweet is enriched with entities before being stored.

WEEK 1 — DAILY WORK LOG (continued)

Day 5 | 22 May 2026 | Data Pipeline, Scraping Execution & Bug Fixes

Day 5 involved completing the storage and graph transformation modules, executing the first full scrape run on all 20 handles, and resolving a series of bugs discovered during the run.

- Built storage/writer.py — atomic JSONL writer for raw data, structured JSON writer for parsed data, session log appender, all-handle merger (merged.json), and scrape report generator.
- Built graph/transformer.py — reads merged.json and produces 16 CSV files: 5 node CSVs (users, tweets, hashtags, links, topics) and 11 relationship CSVs including the new TAGGED_WITH, AUTHORED_ABOUT, and MENTIONED_BY_USER relationships.
- Built graph/cypher_loader.py — connects to Neo4j Desktop, runs constraints.cypher, load_nodes.cypher, and load_relationships.cypher in sequence using the official neo4j Python driver.
- Built main.py — CLI orchestrator with --mode (scrape/transform/load/all), --resume, and --dry-run flags. Executed the full pipeline: python main.py --mode all.
- Bug identified and fixed: display_name selector capturing @handle suffix after newline — fixed with split(newline)[0].strip() in scrape_profile().
- Bug identified and fixed: follower count selector returning 0 for all profiles — replaced with multi-strategy _get_stat_count() trying three selector variations.
- Bug identified and fixed: link extraction returning empty — X injects newlines inside URLs. Fixed with iterative newline collapse regex before URL matching.
- Bug identified and fixed: retweet_of_tweet_id self-referencing — fixed _detect_retweet() to filter out own-handle permalink and match against other-handle status links only.
- Scrape run completed: 20 handles, 1,601 tweets, ~90 minutes runtime with rate limiting.

Day 6 | 23 May 2026 | Neo4j Graph Loading & Analysis

Day 6 focused on loading the transformed CSV data into Neo4j Desktop, verifying the graph structure, and running the analysis and campaign detection Cypher queries.

- Ran python main.py --mode transform to generate all 16 CSVs from merged.json.
- Ran python main.py --mode load to execute constraints.cypher, load_nodes.cypher, and load_relationships.cypher against the local Neo4j Desktop database.
- Verified graph load: 3,215 total nodes (469 users, 1,601 tweets, 303 hashtags, 425 links, 12 topics) and 6,219 relationships across 11 relationship types.
- Executed analysis_queries.cypher sections in Neo4j Browser: top hashtags, most mentioned users, shared hashtag pairs, influence score (retweet count), direct mention network.
- Executed campaign_detection.cypher: hashtag burst detection, hashtag campaign timeline.
- Identified key OSINT findings: Republic TV hashtag dominance, #JusticeForTwisha coordinated campaign, rhythms22 as anomalous high-mention target, narendramodi as broadcast-only handle with no cross-network interaction.
- Cleared self-referencing REPLY_TO relationships in Neo4j Browser: MATCH (t:Tweet)-[r:REPLY_TO]->(t) DELETE r.

Day 7 | 24–25 May 2026 | Analysis, Reporting & Documentation

The final day was dedicated to completing the graph analysis, compiling findings, and preparing the weekly progress report and project documentation.

- Completed all runnable analysis queries from analysis_queries.cypher and campaign_detection.cypher. Documented which queries returned data and which had insufficient data (reply chains, bridge users — insufficient due to scrape window).
- Captured all query result screenshots for inclusion in the weekly report.
- Analysed full graph visualisation in Neo4j Browser — identified per-handle tweet clusters, cross-cluster mention edges, and two isolated handles (RahulGandhi, AltNewsIndia) with no direct cross-network interactions.
- Compiled the complete scrape report: 20 handles scraped, 1,601 tweets, 18 handles with full profile data, 2 partially scraped.
- Written README.md documentation covering installation, configuration, pipeline execution, Neo4j schema, anti-ban design, and extension guide.
- Prepared the Week 1 progress report with full daily breakdown, member role contributions, and analysis findings.

WEEK 1 — MEMBER CONTRIBUTIONS

Rishav Pal — Data Engineer

- Led the design and implementation of the full data pipeline architecture for the linkmap project.
- Built scraper/browser.py and scraper/auth.py — Playwright-based persistent browser session, X/Twitter login automation, cookie persistence, and stealth anti-detection patches.
- Built scraper/rate_limiter.py — anti-ban rate limiting with human scroll simulation, random delay injection, and session health monitoring.
- Implemented and debugged scraper/tweet_scraper.py — profile metadata extraction, timeline scrolling, tweet entity extraction using data-testid selectors.
- Resolved all critical scraper bugs: display_name newline suffix, follower count selector failure, URL newline injection breaking link regex, and retweet self-referencing.
- Built storage/writer.py and storage/schema.py — Pydantic v2 data models with soft validation, atomic JSONL writer, parsed JSON output, session logging, and all-handle merger.
- Configured config/settings.json, config/handles.json (20 India handles), and config/topic_map.json with 12 political/social topic categories.
- Executed the full 20-handle scrape run: 1,601 tweets collected across ~90 minutes.

Sayantana Mandal — OSINT Analyst

- Led the OSINT methodology research and practical intelligence collection exercises on Days 1–2.
- Performed Google Dorking on a target individual using advanced search operators (site:, filetype:, inurl:, intitle:, intext:) to extract publicly available digital footprint data.
- Researched and defined the 20 India-based Twitter/X handles for scraping, selecting handles across political parties, journalism, fact-checking, and civil society for maximum graph richness.
- Designed config/topic_map.json keyword-to-topic mappings with 12 topic categories covering Indian political and social discourse. Iteratively refined keywords to reduce false-positive matches.
- Conducted all Neo4j Browser analysis sessions — executed all sections of analysis_queries.cypher and campaign_detection.cypher, captured screenshots, and documented findings.
- Identified key OSINT intelligence findings: Republic TV's coordinated hashtag campaign, #JusticeForTwisha burst event on 2026-05-22T15, rhythms22 anomalous mention pattern, and RahulGandhi's broadcast-only engagement strategy.
- Analysed the full graph visualisation and interpreted cluster structure, isolation patterns, and cross-network bridge nodes.
- Compiled all analysis screenshots and wrote the findings section of the weekly report.

Shivam Kumar — Graph Analyst

- Led the Neo4j graph database design, schema definition, and Cypher query development.
- Studied Neo4j Desktop, Cypher language, and Neo4j Bloom during Days 1–2. Practised writing MATCH, MERGE, CREATE, and constraint Cypher queries.
- Designed the complete graph schema: 5 node types (User, Tweet, Hashtag, Link, Topic) and 11 relationship types (POSTED, MENTIONED_IN, HAS_HASHTAG, CONTAINS_LINK, REPLY_TO, RETWEET_OF, QUOTES, BELONGS_TO, TAGGED_WITH, AUTHORED_ABOUT, MENTIONED_BY_USER).
- Built graph/transformer.py — transforms merged.json into 16 Neo4j-importable CSV files with correct ID namespacing, domain validation filters, and relationship deduplication.
- Added three new relationship types to improve graph richness: TAGGED_WITH (Tweet→Topic direct), AUTHORED_ABOUT (User→Topic with frequency threshold ≥ 2), and MENTIONED_BY_USER (User→User with count).
- Wrote all Cypher files: constraints.cypher (uniqueness + indexes), load_nodes.cypher, load_relationships.cypher, analysis_queries.cypher (15 sections), and campaign_detection.cypher (5 sections).
- Built graph/cypher_loader.py — automated Neo4j pipeline loader with sequence control, error handling for duplicate constraints, and row-count logging.
- Verified graph integrity post-load: identified and cleaned self-referencing REPLY_TO edges. Confirmed 3,215 nodes and 6,219 relationships successfully loaded.
- Interpreted graph topology: identified journalism cluster as most central hub, confirmed RahulGandhi isolation, and flagged republic→rhythms22 (49 mentions) as anomaly.

WEEK 1 — GRAPH ANALYSIS RESULTS

The following analysis was performed using Neo4j Browser by executing queries from analysis_queries.cypher and campaign_detection.cypher against the fully loaded graph containing 3,215 nodes and 6,219 relationships. Screenshots of query results are provided below.

Graph Statistics Summary

Entity	Count	Relationship	Count
Users	469	POSTED	2,012
Tweets	1,601 scraped	MENTIONED_IN	1,095
Hashtags	303	HAS_HASHTAG	1,833
Links	425	CONTAINS_LINK	542
Topics	12	TAGGED_WITH	859
		AUTHORED_ABOUT	93
		MENTIONED_BY_USER	437
		RETWEET_OF	300

Full Graph Visualisation — All Nodes & Relationships

The graph below shows all scraped handles as central dark nodes, surrounded by their tweet clouds (pink) and hashtag/link nodes (cyan). Cross-cluster connecting lines represent MENTIONED_BY_USER relationships — the inter-handle influence network. Two clusters (RahulGandhi, AltNewsIndia) are isolated with no cross-network mention edges, indicating broadcast-only communication strategies.




Figure 1: Full graph — 3,215 nodes (pink=Tweets, cyan=Hashtags/Links, purple=Users)

WEEK 1 — QUERY RESULTS & FINDINGS

Finding 1: Top 10 Hashtags by Frequency

The top hashtags are dominated entirely by Republic TV branded tags: #RepublicTV (92), #RepublicExclusive (89), #ArnabGoswami (89), #RepublicDigital (89). This indicates that Republic TV uses a highly systematic hashtag strategy — virtually every tweet includes their branded hashtags. #JusticeForTwisha (78) appearing at rank 7 is a significant finding as it represents a coordinated campaign hashtag (confirmed by burst detection). #Breaking (43) and #Watch (42) are generic broadcast tags shared across media handles.

h.tag	uses
1 "republictv"	92
2 "republicexclusive"	89
3 "arnabgoswami"	89
4 "republicdigital"	89
5 "livenews"	83
6 "indianews"	83
7 "justicefortwisha"	78
8 "republicworldlive"	55
9 "breaking"	43
10 "watch"	42

Figure 2: Top 10 hashtags by frequency across all 20 handles

Finding 2: Top 10 Most Mentioned Users

rhythms22 (55 mentions) ranks #1 despite being outside the 20 scraped handles — all 55 mentions originate from Republic TV alone, which is a strong anomaly signal. narendramodi (38) as #2 reflects his central political position — mentioned by BJP-aligned handles and journalists. swedishpm (23) at #3 reflects a diplomatic event during the scrape window. shekhargupta (22) being heavily mentioned by ThePrintIndia reflects institutional cross-promotion. adityarajkaul (16) and gargirawat (11) are NDTV correspondents consistently tagged in their organisation's tweets.

u.handle	mentions
1 "rhythms22"	55
2 "narendramodi"	38
3 "swedishpm"	23
4 "shekhargupta"	22
5 "vedikas"	17
6 "adityarajkaul"	16
7 "svembu"	14
8 "ashwinivaishnaw"	13
9 "anurag_dwary"	13
10 "gargirawat"	11

Figure 3: Top 10 most mentioned users across all 20 handles

WEEK 1 — QUERY RESULTS & FINDINGS (continued)

Finding 3: Users Sharing the Same Hashtag (Coordination Signals)

ndtv ↔ republic sharing #Breaking 408 times is the single strongest co-occurrence pair, indicating both organisations cover breaking news using the same tag — a content competition signal rather than coordination. ShekharGupta ↔ ThePrintIndia dominate 6 entries in the top 14 pairs (#ThePrintOpinion 80, #ThePrintFeature 52, #CutTheClutter 44, #ThePrintPolitics 36, #ThePrintGroundReport 28, #NationalInterest 24) — confirming strong institutional alignment between the founder and the publication. bdutt ↔ republic sharing #JusticeForTwisha (77) and #TwishaSharma (37) suggests cross-organisation coverage of the same campaign event.

u1.handle	u2.handle	h.tag	shared_hashtag_count
"ndtv"	"republic"	"breaking"	408
"republic"	"sardesaiarajdeep"	"indianews"	82
"ShekharGupta"	"ThePrintIndia"	"theprintopinion"	80
"ndtv"	"smitaprakash"	"watch"	78
"bdutt"	"republic"	"justicefortwisha"	77
"ShekharGupta"	"ThePrintIndia"	"theprintfeature"	52
"ShekharGupta"	"ThePrintIndia"	"cuttheclutter"	44
"sardesaiarajdeep"	"smitaprakash"	"watch"	39
"bdutt"	"republic"	"twishasharma"	37
"republic"	"sardesaiarajdeep"	"twishasharma"	37
"ShekharGupta"	"ThePrintIndia"	"theprintpolitics"	36
"ShekharGupta"	"ThePrintIndia"	"theprintgroundreport"	28
"ShekharGupta"	"ThePrintIndia"	"nationalinterest"	24
"Zoo_bear"	"ndtv"	"breaking"	24

Figure 4: Top user pairs sharing the same hashtag

Finding 4: Influence Score — Retweets Received

smitaprakash leads with 84 retweets received, followed by Zoo_bear (Mohammed Zubair) at 56 and sardesaiarajdeep at 43. This is notable because smitaprakash and Zoo_bear are not the highest-follower accounts — their high retweet scores indicate their content generates disproportionate amplification relative to account size, making them high-influence nodes in the network. Tejasvi_Surya (17) and internetfreedom (19) appear above ndtv (17) and dhruv_rathee (16), suggesting niche but highly engaged audiences for political and policy content.

u.handle	retweets_received
"smitaprakash"	84
"Zoo_bear"	56
"sardesaiarajdeep"	43
"bdutt"	35
"internetfreedom"	19
"Tejasvi_Surya"	17
"ndtv"	17
"dhruv_rathee"	16
"ShekharGupta"	5
"priyankagandhi"	5

Figure 5: Influence score — retweets received per handle

WEEK 1 — CAMPAIGN DETECTION & MENTION NETWORK

Finding 5: Direct Mention Network (User → User)

The direct mention network reveals how scraped handles interact with each other and external accounts. ndtv → adityarajkaul (82) and ndtv → gauried (54) are the two strongest edges — NDTV consistently tags its own correspondents in stories, a standard broadcast practice. republic → rhythms22 (49) remains the most anomalous edge — Republic TV mentioning an unverified external account 49 times warrants further OSINT investigation. Tejasvi_Surya → narendramodi (30) confirms the BJP MP's amplification strategy for the PM. ndtv → secrubio (15) reflects a specific diplomatic event (US Secretary of State visit) during the scrape window. ThePrintIndia → shekhargupta (15) confirms institutional cross-promotion.

mentioner	mentioned	times
"ndtv"	"adityarajkaul"	82
"ndtv"	"gauried"	54
"republic"	"rhythms22"	49
"Tejasvi_Surya"	"narendramodi"	30
"ndtv"	"vishnundtv"	24
"narendramodi"	"swedishpm"	19
"ndtv"	"tejshreethought"	18
"ndtv"	"aishvaryjain"	18
"ThePrintIndia"	"shekhargupta"	15
"ndtv"	"secrubio"	15
"republic"	"svembu"	14
"ndtv"	"dharanvija47684"	12
"ndtv"	"anurag_dwary"	11
"ndtv"	"anushkagarg2000"	11

Figure 6: Direct mention network — top mention pairs with frequency

Finding 6: Hashtag Co-occurrence Burst Detection (Campaign Detection)

The burst detection query identified a coordinated hashtag campaign on 2026-05-22 at 15:00 UTC (8:30 PM IST). The hashtag cluster {#ArnabGoswami, #Giribala, #GiribalaCoverup, #IndiaNews, #JusticeForTwisha, #LiveNews, #RepublicWorldLive} all co-occurred in 14 tweets within the same one-hour window. The identical co-occurrence count (14) across all pairs in the cluster confirms this was a single coordinated posting event — all 14 tweets used the full set of hashtags simultaneously. This is a textbook coordinated hashtag campaign pattern, with Republic TV as the originating handle.

tag1	tag2	hour_bucket	co_occurrences
"arnabgoswami"	"giribala"	"2026-05-22T15"	14
"arnabgoswami"	"giribalacoverup"	"2026-05-22T15"	14
"giribala"	"giribalacoverup"	"2026-05-22T15"	14
"arnabgoswami"	"indianews"	"2026-05-22T15"	14
"giribalacoverup"	"indianews"	"2026-05-22T15"	14
"giribala"	"indianews"	"2026-05-22T15"	14
"indianews"	"justicefortwisha"	"2026-05-22T15"	14
"arnabgoswami"	"justicefortwisha"	"2026-05-22T15"	14
"giribalacoverup"	"justicefortwisha"	"2026-05-22T15"	14
"giribala"	"justicefortwisha"	"2026-05-22T15"	14
"justicefortwisha"	"livenews"	"2026-05-22T15"	14
"indianews"	"livenews"	"2026-05-22T15"	14
"arnabgoswami"	"livenews"	"2026-05-22T15"	14
"giribalacoverup"	"livenews"	"2026-05-22T15"	14

Figure 7: Hashtag co-occurrence burst — 2026-05-22T15:00 UTC

WEEK 1 — ADDITIONAL FINDINGS & NEXT STEPS

Finding 7: Hashtag Campaign Timeline

The hashtag timeline query for the #TheRealTraitor campaign shows a single tweet from Republic TV on 2026-05-20 at 16:17 UTC, with 1 retweet. The tweet text frames a debate segment pitting Congress and BJP spokespersons against each other, with the campaign hashtag designed to generate partisan engagement. While this specific instance shows limited spread (1 RT), the framing — using a provocative hashtag to introduce a televised political debate — is a recurring tactic in the Republic TV dataset.

time	author	tweet_text	retweets
"2026-05-20T16:17:31+00:00"	"republic"	#TheRealTraitor 'PM Narendra Modi is not the centre of my universe': Faizan Kidwai, National Spokesperson, SP(@FaizanFKidwai) Versus Shazia Ilmi, National Spokesperson, BJP(@shaziailmi) Watch #DebateWithArbab now, on-air, and online. Tune in and fire in your comments: https://youtube.com/live/neGxSuFR-Qo?si=Wavf6LQh4nZXEqTJ #PMModi #RahulGandhi #BJP #Congress #RepublicTV #IndiaNews #Live	1

Figure 8: Hashtag campaign timeline — #TheRealTraitor

Finding 8: Specific Link Domain Analysis

The specific link query shows ThePrint as the dominant content source in the dataset, with ShekharGupta and ThePrintIndia consistently sharing article links to theprint.in. Article topics cover governance (Delhi Gymkhana land dispute), judicial decisions (school fee regulation), military/social issues (Army wife story), and regional politics (Ladakh talks). All links are accompanied by correspondent bylines, confirming ThePrint's practice of tagging reporters in every shared article — a pattern that generates both the MENTIONED_BY_USER edges and the link nodes in the graph simultaneously.

tweet	Expanded_url
"Now this is audacity. Modi govt takes on the real and permanent establishment of India..."	"https://t.co/39fAMf8K1a"
Centre tells Delhi Gymkhana Club to hand over land by 5 June, needs it to 'secure defence infra'	
Moushumi Das Gupta @dgupta_moushumi reports for ThePrint	
"At least the judges have the good sense to block politico-bureaucratic overkill."	"https://t.co/TM3wW0A3P"
In battle over fee hike regulation, Delhi HC rules in favour of private schools, limits govt's power	
Alfreza Akwad, TPSJ alum & intern at ThePrint, reports	
"I am an Army wife, and I feel a personal debt to Pune's Artificial Limb Centre"	"https://t.co/1D4xw8K1b"
Anubrem Zaidi @AnubremZaidi, award-winning military author and activist, writes #ThePrintOpinion	
"Breakthrough in Ladakh talks, Centre & local groups agree on 'elected representative' body, safeguards	"https://t.co/3J2b45Y8KES"

Figure 9: Link domain analysis — ThePrint article links

Queries Without Sufficient Data (This Week)

The following queries from analysis_queries.cypher and campaign_detection.cypher returned no results due to insufficient data within the current scrape window. These will be revisited after expanding the handle set in Week 2.

- Reply chain analysis — only 3 valid REPLY_TO edges after self-reference cleanup (insufficient for path traversal)
- Bridge user detection — no users found spanning 2+ topic clusters via mention edges in current dataset
- Most retweeted narrative by topic — RETWEET_OF has 300 edges but cross-topic tweet matching insufficient for per-topic ranking
- Topic cluster visualisation — AUTHORED_ABOUT has 93 edges but community detection needs more inter-handle connections